

# Homework 2

Shen Dingtao 3170104764

## 1. Loading and cleaning

(a)

```
ca_pa <- read.csv("data/calif_penn_2011.csv",header=T)
```

(b)

```
dim(ca_pa)
```

```
## [1] 11275    34
```

So the data frame has 11275 rows and 34 columns.

(c)

```
colSums(apply(ca_pa,c(1,2),is.na))
```

```
##                X                GEO.id2
##                0                0
##            STATEFP            COUNTYFP
##                0                0
##            TRACTCE            POPULATION
##                0                0
##            LATITUDE            LONGITUDE
##                0                0
##    GEO.display.label    Median_house_value
##                0                599
##            Total_units            Vacant_units
##                0                0
##            Median_rooms    Mean_household_size_owners
##                157                215
##    Mean_household_size_renters    Built_2005_or_later
##                152                98
##            Built_2000_to_2004    Built_1990s
##                98                98
##            Built_1980s            Built_1970s
##                98                98
##            Built_1960s            Built_1950s
##                98                98
##            Built_1940s    Built_1939_or_earlier
##                98                98
##            Bedrooms_0            Bedrooms_1
```

```
##           98           98
##           Bedrooms_2       Bedrooms_3
##           98           98
##           Bedrooms_4       Bedrooms_5_or_more
##           98           98
##           Owners           Renters
##           100           100
## Median_household_income Mean_household_income
##           115           126
```

This command is used to count the number of NA (not a number) values in each column of the data frame.

- The `apply()` function loops through all the elements of the matrix `ca_pa` and applies the `is.na()` function which returns TRUE if the element is not a number and FALSE otherwise.
- The resulting matrix of TRUE and FALSE values is then given as input to `colSums()` function, which counts the number of TRUE values in each column.

(d)

```
new_ca_pa<-na.omit(ca_pa)
```

(e)

```
nrow(ca_pa)-nrow(new_ca_pa)
```

```
## [1] 670
```

The result shows that the `na.omit()` command eliminate 670 rows.

(f) My answer in (c) and (e) are compatible. We can run the following command to verify this:

```
colSums(apply(new_ca_pa,c(1,2),is.na))
```

```
##           X           GEO.id2
##           0           0
##           STATEFP       COUNTYFP
##           0           0
##           TRACTCE       POPULATION
##           0           0
##           LATITUDE      LONGITUDE
##           0           0
## GEO.display.label      Median_house_value
##           0           0
##           Total_units   Vacant_units
##           0           0
##           Median_rooms  Mean_household_size_owners
##           0           0
## Mean_household_size_renters      Built_2005_or_later
##           0           0
##           Built_2000_to_2004      Built_1990s
##           0           0
```

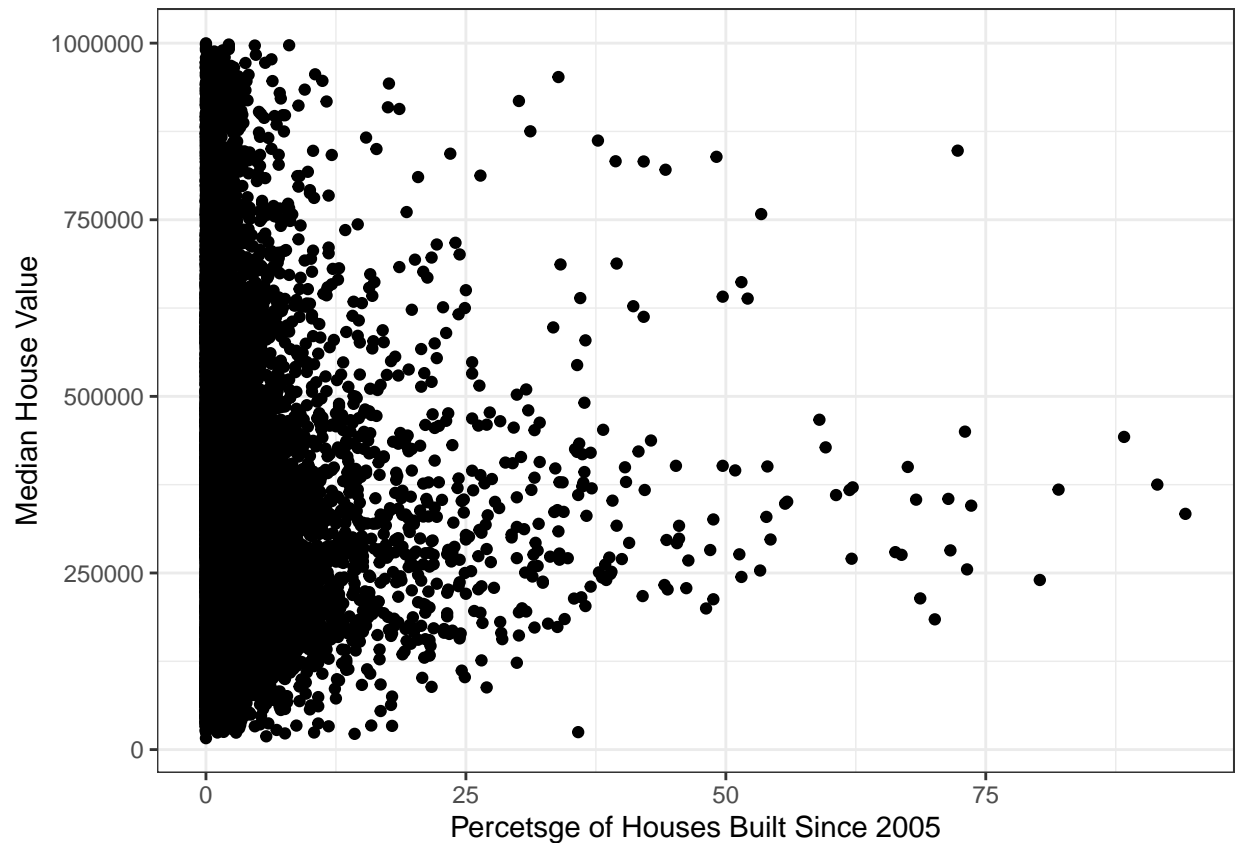
##	Built_1980s	Built_1970s
##	0	0
##	Built_1960s	Built_1950s
##	0	0
##	Built_1940s	Built_1939_or_earlier
##	0	0
##	Bedrooms_0	Bedrooms_1
##	0	0
##	Bedrooms_2	Bedrooms_3
##	0	0
##	Bedrooms_4	Bedrooms_5_or_more
##	0	0
##	Owners	Renters
##	0	0
##	Median_household_income	Mean_household_income
##	0	0

The values of all columns are all 0, which means that we have successfully purged `ca_pa` from any row containing NA value.

## 2. *This Very New House*

(a)

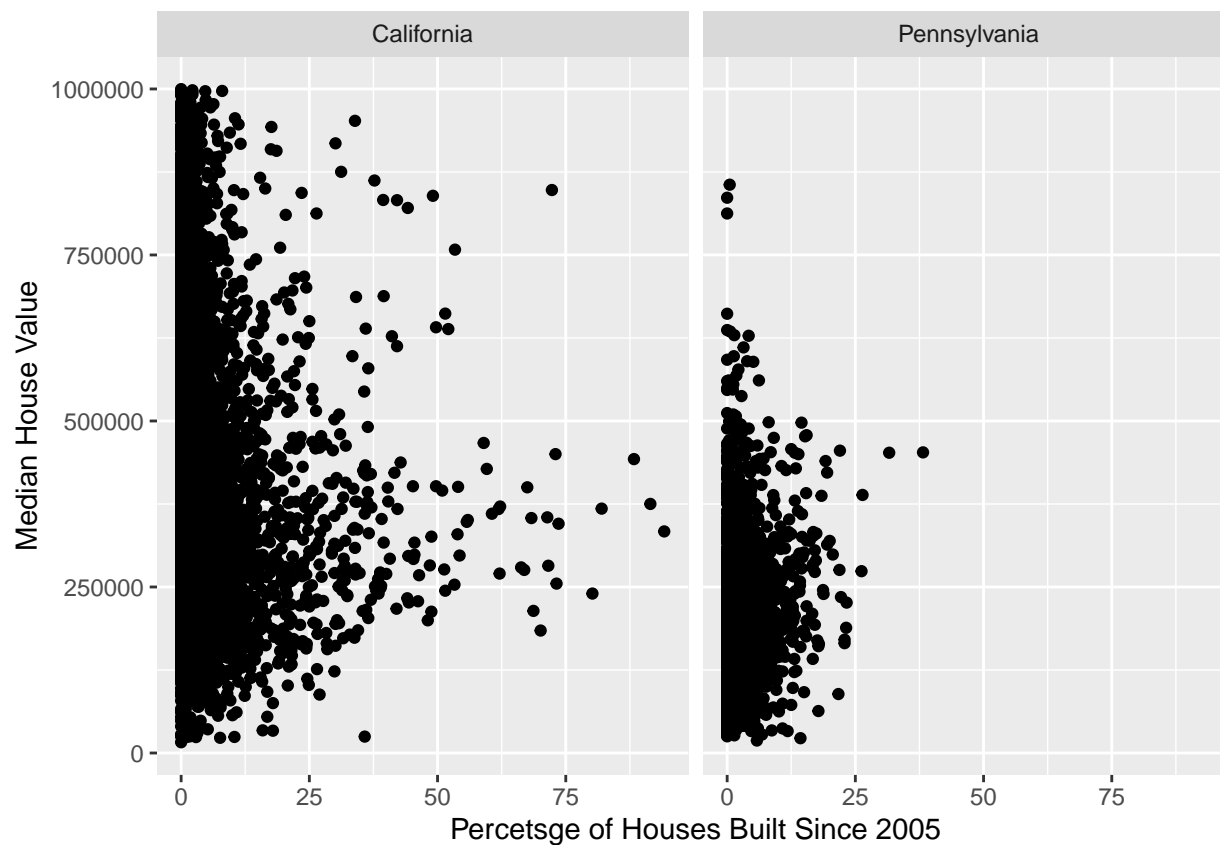
```
ggplot(data = new_ca_pa) +
  geom_point(aes(x = Built_2005_or_later,
    y = Median_house_value), na.rm = TRUE) +
  labs(x = "Percentage of Houses Built Since 2005",
    y = "Median House Value") +
  theme_bw() +
  theme(legend.title=element_blank())
```



(b) With following commands, we make a pair of plots that breaks this out by state, which shows the median house prices against Built\_2005\_or\_later in California(state 6) and Pennsylvania(state 42) respectively.

```
d<-data.frame(STATEFP=c(6,42),state=c("California","Pennsylvania"))
ca_pa2<-left_join(new_ca_pa,d,by="STATEFP")
```

```
ggplot(data = ca_pa2) +
  geom_point(aes(x = Built_2005_or_later,
    y = Median_house_value),na.rm = TRUE) +
  labs(x = "Percetsge of Houses Built Since 2005",
    y = "Median House Value") +facet_wrap(~ state)
```



### 3. *Nobody Home*

(a) Add a new column to the dataframe which contains the vacancy rate.

```
ca_pa3 <- ca_pa2 %>%  
  mutate(Vacant_Rate=Vacant_units/Total_units)
```

Minimum:

```
min(ca_pa3$Vacant_Rate,na.rm = TRUE)
```

```
## [1] 0
```

Maximum:

```
max(ca_pa3$Vacant_Rate,na.rm = TRUE)
```

```
## [1] 0.965311
```

Mean:

```
mean(ca_pa3$Vacant_Rate,na.rm = TRUE)
```

```
## [1] 0.08888789
```

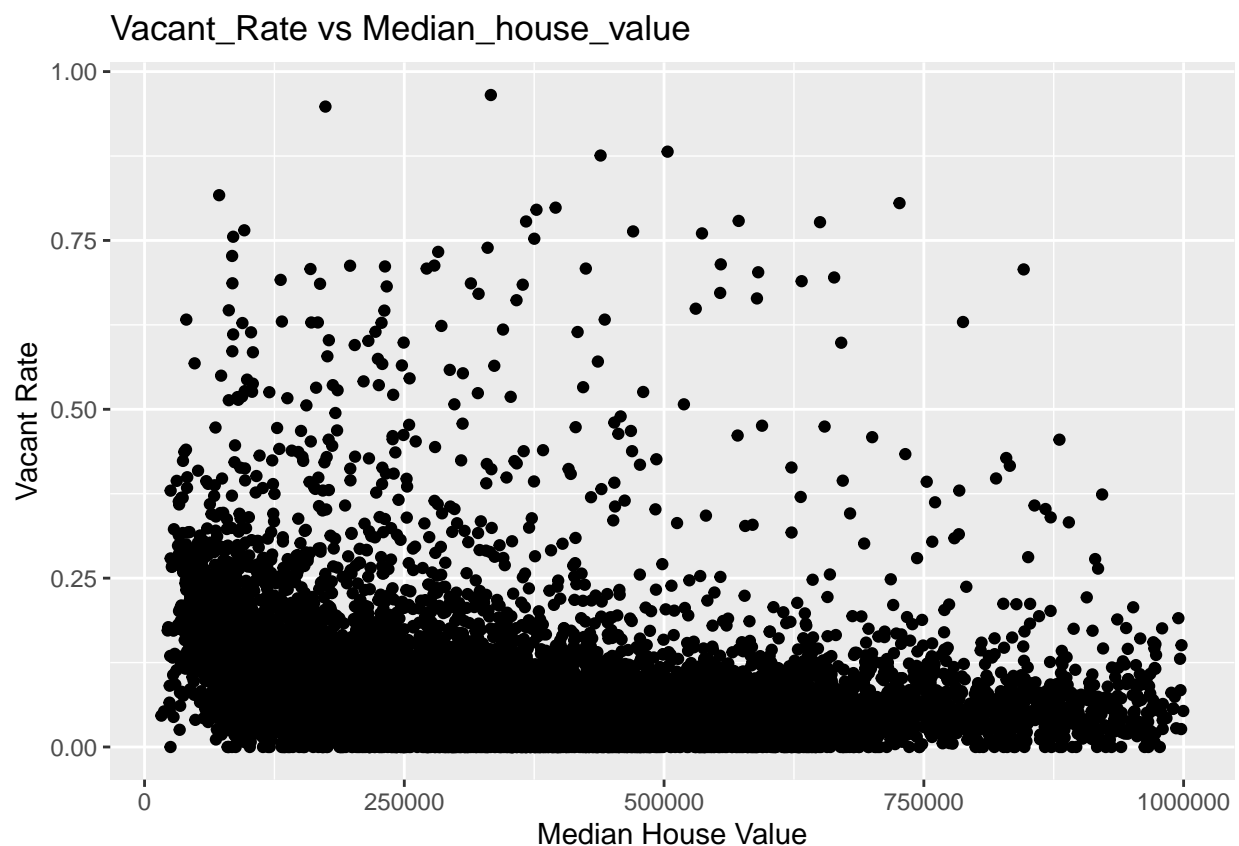
Median:

```
median(ca_pa3$Vacant_Rate,na.rm = TRUE)
```

```
## [1] 0.06767283
```

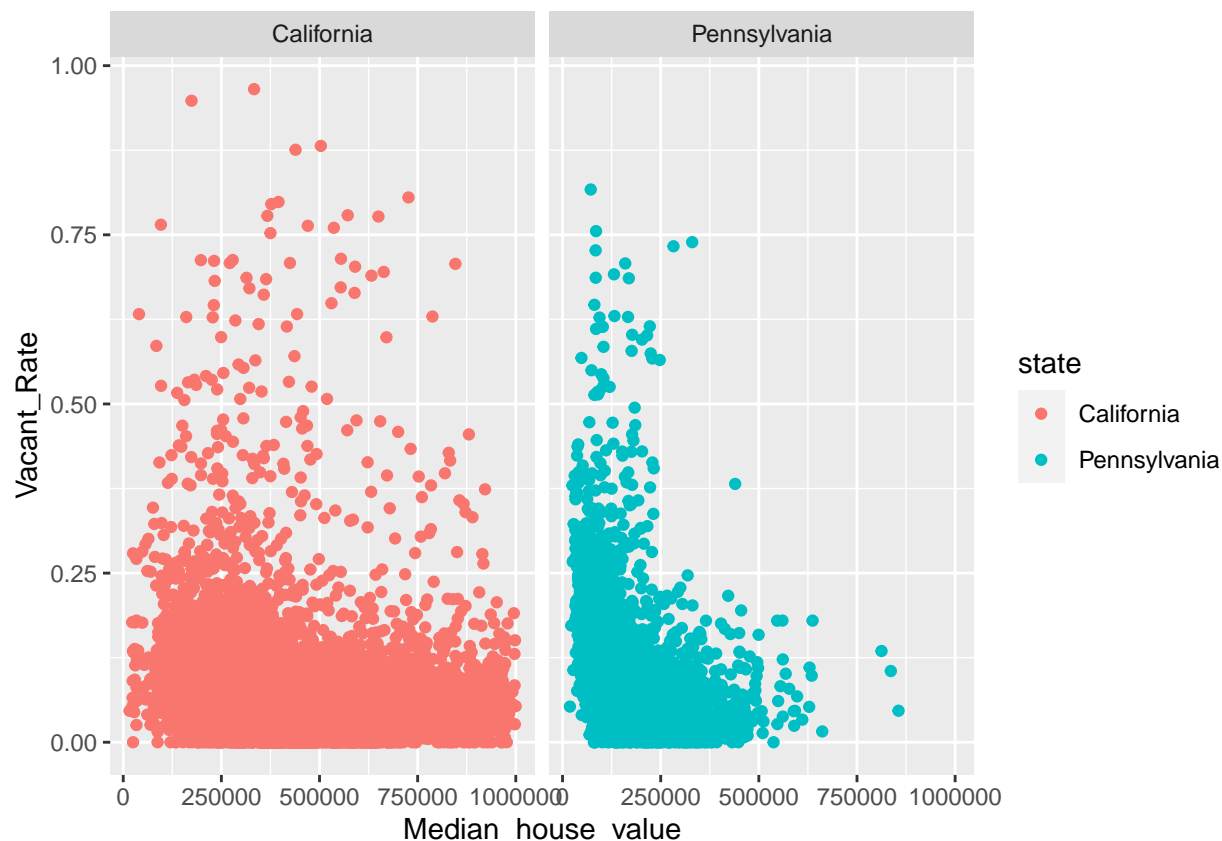
(b) Plot the vacancy rate against median house value:

```
ggplot(data = ca_pa3) +  
  geom_point(aes(x = Median_house_value,  
    y = Vacant_Rate),na.rm = TRUE) +  
  labs(x = "Median House Value",  
    y = "Vacant Rate",title = "Vacant_Rate vs Median_house_value")
```



(c) Plot vacancy rate against median house value separately for California and for Pennsylvania:

```
ggplot(data = ca_pa3) +  
  geom_point(aes(x = Median_house_value,  
    y = Vacant_Rate,color=state),na.rm = TRUE) +  
  facet_wrap(~ state)
```



4.

(a) The block of code is supposed to calculate the median house value in Alameda County (country 1 in California). It firstly selects the rows of California, whose STATEFP=6, and then selects the rows of Alameda, whose COUNTYFP=1 among the selected rows. Finally select median value of these rows' Median house value to get the median house value in Alameda.

(b) We can obtain the same result as the block of code with following command:

```
ca_pa3 %>% filter(STATEFP==6,COUNTYFP==1) %>% {median(.$Median_house_value,na.rm = TRUE)}
```

```
## [1] 474050
```

(c) We can obtain the average percentages of housing built since 2005 for Alameda, Santa Clara and Alameda Counties with following commands: (i) For Alameda:

```
(Alameda_avg <- ca_pa3 %>% filter(STATEFP==6&COUNTYFP==1 )) %>% {mean(.$Built_2005_or_later,na.rm = TRUE)}
```

```
## [1] 2.820468
```

(ii) For Santa Clara:

```
(Alameda_avg <- ca_pa3 %>% filter(STATEFP==6&COUNTYFP==85 )) %>% {mean(.$Built_2005_or_later,na.rm = TRUE)}
```

```
## [1] 3.200319
```

(ii) For Allegheny:

```
(Alameda_avg <- ca_pa3 %>% filter(STATEFP==42&COUNTYFP==3 )) %>% {mean(.$Built_2005_or_later,na.rm = TR
```

```
## [1] 1.474219
```

(d) the correlation between median house value and the percent of housing built since 2005 (i) In the whole data:

```
cor(ca_pa3$Median_house_value,ca_pa3$Built_2005_or_later)
```

```
## [1] -0.01893186
```

(ii) In all of California

```
ca_pa3 %>% filter(STATEFP==6) %>%  
{cor(.$Median_house_value,.$Built_2005_or_later)}
```

```
## [1] -0.1153604
```

(iii) In all of Pennsylvania

```
ca_pa3 %>% filter(STATEFP==42) %>%  
{cor(.$Median_house_value,.$Built_2005_or_later)}
```

```
## [1] 0.2681654
```

(iv) In Alameda County

```
ca_pa3 %>% filter(STATEFP==6,COUNTYFP==1) %>%  
{cor(.$Median_house_value,.$Built_2005_or_later)}
```

```
## [1] 0.01303543
```

(v) In Santa Clara County

```
ca_pa3 %>% filter(STATEFP==6,COUNTYFP==85) %>%  
{cor(.$Median_house_value,.$Built_2005_or_later)}
```

```
## [1] -0.1726203
```

(vi) In Allegheny Count

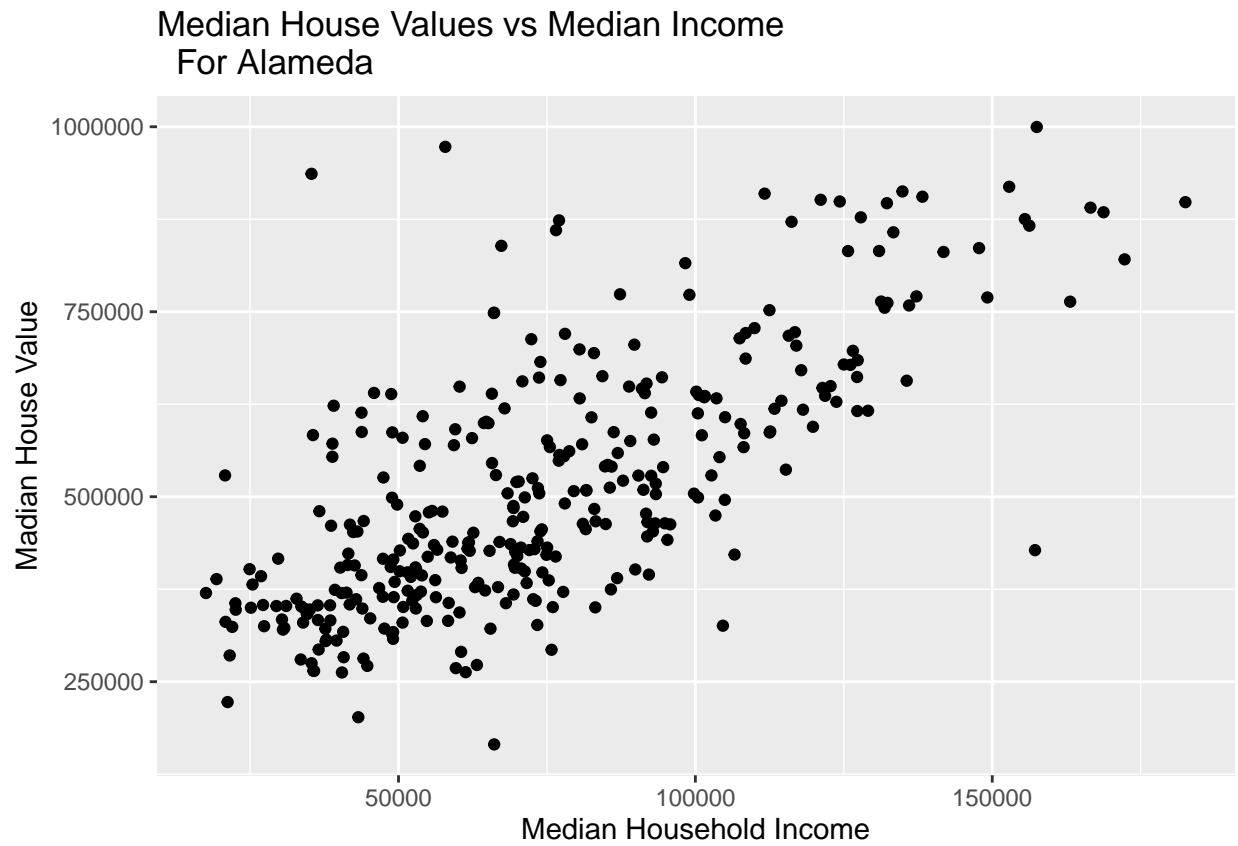
```
ca_pa3 %>% filter(STATEFP==42,COUNTYFP==3) %>%  
{cor(.$Median_house_value,.$Built_2005_or_later)}
```

```
## [1] 0.1939652
```

(e) Median house values against median income. (i) For Alameda:




```
ca_pa3 %>% filter(STATEFP==6,COUNTYFP==1) %>%
  ggplot() +
  geom_point(aes(x = Median_household_income,
    y = Median_house_value),na.rm = TRUE) +
  labs(x = "Median Household Income",
    y = "Madian House Value",title = "Median House Values vs Median Income
    For Alameda")
```



(ii) For Santa Clara:


```
ca_pa3 %>% filter(STATEFP==6,COUNTYFP==85) %>%
  ggplot() +
  geom_point(aes(x = Median_household_income,
    y = Median_house_value),na.rm = TRUE) +
  labs(x = "Median Household Income",
    y = "Madian House Value",title = "Median House Values vs Median Income
    For Santa Clara")
```



(iii) For Allegheny:

```
ca_pa3 %>% filter(STATEFP==42,COUNTYFP==3) %>%
  ggplot() +
  geom_point(aes(x = Median_household_income,
    y = Median_house_value),na.rm = TRUE) +
  labs(x = "Median Household Income",
    y = "Median House Value",title = "Median House Values vs Median Income
    For Allegheny")
```

## Median House Values vs Median Income For Allegheny



MB.Ch1.11. Run the following code:

```
gender <- factor(c(rep("female", 91), rep("male", 92)))
table(gender)
```

```
## gender
## female  male
##      91    92
```

```
gender <- factor(gender, levels=c("male", "female"))
table(gender)
```

```
## gender
##  male female
##    92     91
```

```
gender <- factor(gender, levels=c("Male", "female"))
# Note the mistake: "Male" should be "male"
table(gender)
```

```
## gender
##  Male female
##    0     91
```

```
table(gender, exclude=NULL)
```

```
## gender
##   Male female  <NA>
##     0     91    92
```

```
rm(gender) # Remove gender
```

**Explain:** `table` uses the cross-classifying factors to build a contingency table of the counts at each combination of factor levels. That is, it is used to calculate frequency.

- i) In the first command, `factor` turns a vector into a factor data type. And then `gender` has 91 'female's and 92 'male's, with levels ("female", "male").
- ii) In the second command, attribute `levels=()` is used to return the value of the levels of its argument, which select values that fit levels in the data as valid levels. Then we use `table` to count frequency of the levels ("male", "female"), in order.
- iii) In the third command, attribute `levels=()` is used to set the attribute, which makes "Male" that is not in `gender` a level, and the number of its value is 0, and all "male"s in `gender` is discarded as invalid values.
- iv) In the fourth command, we use `exclude=NULL` in `table()` to count all the invalid values. So previous discarded "male"s are counted and was outputted as :92

## MB.Ch1.12.

```
k<-0
Func <- function(x,cutoff_value){
  for(i in x){
    if (i>cutoff_value){
      k<-k+1
    }
  }
  prop <- k/length(x)
  return(prop)
}
```

- (a) For the sequence of numbers 1, 2, . . . , 100, we set the value cutoff 35, 67.3 and 89, then the expected proportion result is 0.65, 0.33 and 0.11 correspondingly. And following codes check this.

```
x<-1:100
vc<-35
Func(x,vc)
```

```
## [1] 0.65
```

```
vc<-67.3
Func(x,vc)
```

```
## [1] 0.33
```

```
vc<-89
Func(x,vc)
```

```
## [1] 0.11
```

**MB.Ch1.18.** Using following commands, we can convert Rabbit to the required form.

```
Dose <- unstack(Rabbit, Dose ~ Animal)[,1]
Treatment <- unstack(Rabbit, Treatment ~ Animal)[,1]
BPchange <- unstack(Rabbit, BPchange ~ Animal)
Rabbit.df <- data.frame(Treatment, Dose, BPchange)
Rabbit.df
```

```
##      Treatment   Dose   R1    R2    R3    R4    R5
## 1    Control    6.25  0.50   1.00  0.75  1.25  1.5
## 2    Control   12.50  4.50   1.25  3.00  1.50  1.5
## 3    Control   25.00 10.00   4.00  3.00  6.00  5.0
## 4    Control   50.00 26.00  12.00 14.00 19.00 16.0
## 5    Control  100.00 37.00  27.00 22.00 33.00 20.0
## 6    Control  200.00 32.00  29.00 24.00 33.00 18.0
## 7         MDL    6.25  1.25   1.40  0.75  2.60  2.4
## 8         MDL   12.50  0.75   1.70  2.30  1.20  2.5
## 9         MDL   25.00  4.00   1.00  3.00  2.00  1.5
## 10        MDL   50.00  9.00   2.00  5.00  3.00  2.0
## 11        MDL  100.00 25.00  15.00 26.00 11.00  9.0
## 12        MDL  200.00 37.00  28.00 25.00 22.00 19.0
```